

# **A Project Report on STOCK DATA ANALYSIS & PREDICTION**

Submitted by

**Sandeep Reddy Panyala – AP24122060005**

**Sri Harsha Vardhan Gogisetty – AP24122060007**

**Eswar Maddi – AP24122060014**



**DSC 502L**

**BIG DATA ANALYTICS**

**Under the guidance of**

**Dr. Rajiv Senapathi**

**Assistant Professor**

**Department of Computer Science and Engineering**

**SRM University-AP**

**Neerukonda, Mangalagiri, Guntur**

**Andhra Pradesh – 522 240**

**Dec 2024**

## Certificate

Date: 02-Dec-24

This is to certify that the work present in this Project entitled “**Stock Data Analysis and Prediction**” has been carried out by **Sandeep Reddy Panyala, Sri Harsha Vardhan Gogisetty, Eswar Maddi** under my/our supervision. The work is genuine, original, and suitable for submission to the SRM University – AP for the award of **Master of Technology** in School of Engineering and Sciences.

## Supervisor

Dr. Rajiv Senapati  
Assistant Professor  
Department of Computer Science and Engineering.

## TABLE OF CONTENTS

1	Abstract.....	1
2	Introduction .....	2
3	Literature Review.....	3
4	Methodology.....	4
4.1	MapReduce.....	4
4.2	Sqoop.....	7
4.3	Hive .....	11
4.4	Pyspark .....	18
4.4.1	Loading of Data.....	19
4.4.2	Data Preprocessing .....	19
4.4.3	Exploratory Data Analysis.....	20
4.4.4	Model Selection .....	29
4.4.5	Model Evaluation.....	30
4.4.6	Predictions.....	31
5	Conclusion .....	33
6	References .....	34

# 1 Abstract

The Stock market analysis and prediction play a pivotal role in understanding market trends and assisting investors in making informed decisions. This project leverages big data technologies like Apache Hadoop, Hive, Sqoop, and PySpark to process and analyze large-scale stock market data. The implementation focuses on aggregating key statistics such as average and maximum values of stock metrics (Close, Volume, Open, High, Low prices) across monthly and yearly timeframes using the MapReduce programming model.

This project explores a comprehensive dataset comprising 12 distinct stocks over a five-year period from 2018 to 2023. The project begins by importing raw stock market data from relational databases into the Hadoop Distributed File System (HDFS) using Sqoop. Hive is then employed as a central data warehouse for managing and querying the datasets, enabling high-level SQL-like queries for efficient data analysis.

MapReduce is implemented to compute key stock metrics grouped by monthly and yearly periods. The Mapper generates grouping keys and emits metric records. The Reducer aggregates these records to compute averages and track maximum values for each metric.

To enhance processing speed and scalability, PySpark is used to implement the analysis in a distributed computing environment. The PySpark API provides an intuitive interface for handling large datasets, ensuring seamless data transformation and parallel computation.

The resulting system integrates stock market data, performs detailed analytics, and provides actionable insights into stock performance trends. By combining Hadoop's storage capabilities, Hive's querying power, Sqoop's data transfer efficiency, and PySpark's computational strength, this project delivers a robust framework for stock price analysis and prediction. This enables better data-driven decision-making in financial markets.

## 2 Introduction

The stock market, an active marketplace where people and organizations invest in and trade shares, is a fundamental component of current finance. This market provides a platform for wealth generation and diversifying investments in addition to reflecting the economic health of countries. The amount of data produced in the stock market has increased as technological and globalized developments change the financial environment, offering analysts and investors equal opportunities and problems.

The capacity to extract valuable insights from large datasets has become critical in times marked by overload of information. An extremely valuable resource for analyzing market trends and behaviors is historical stock data, which comprises a wide range of measures such as ticker symbols, opening and closing prices, daily highs and lows, and trading volumes. Stakeholders may identify trends, evaluate risk, and make well-informed decisions based on information rather than conjecture by analyzing this data.

This project focuses on an extensive dataset that includes 12 different equities over the course of five years, from 2018 to 2023. Through the use of data processing and analytical frameworks, particularly Hive and PySpark, this study seeks to understand the complexities of stock price fluctuations and the fundamental causes of them. For investigating other aspects of stock performance, such as instability, trading volume, and past price movements, the dataset offers a solid starting point.

Key objectives of this analysis include identifying patterns in stock performance, evaluating the impact of market events, and forecasting future price movements. Through exploratory data analysis (EDA) and machine learning methodologies, this research seeks to provide actionable insights that can enhance investment strategies and contribute to a deeper understanding of market dynamics.

### 3 Literature Review

Using big data analytics and machine learning, Zhihao Peng et al. [1] created a data pipeline based on Cloudera and Hadoop to predict daily growth in U.S. oil stocks. In addition to using real-time stock data from Yahoo Finance, they integrated tools such as Hadoop for distributed storage, Apache Spark for real-time analytics, and Flume for data injection. With an R-squared value of only 0.03 and a Mean Absolute Error (MAE) of 1.95%, the findings of a linear regression model used to predict stock price correlations were subpar, suggesting that linear regression is not appropriate for high-dimensional stock data.

To overcome the difficulties associated with high-dimensional financial data, Jithina Jose et al. [2] provide a strategy for stock prediction and analysis that makes use of Hadoop-based methodologies. The method optimizes trading tactics and improves forecast accuracy by combining logistic regression, genetic algorithms, and linear regression. The authors use a genetic algorithm to evolve the best parameters for trading strategies including Head and Shoulders, Double Bottom, and Double Top by converting NASDAQ order book data into tick data. Simulations conducted in real time demonstrate that these strategies outperform random trading in terms of earnings. The study concludes that an effective strategy for algorithmic trading is to combine evolutionary algorithms with pattern-detection methods.

K V D Kiran et al. [3], discusses a system for analysing historical stock data using Hadoop and Hue. It focuses on the impact of macroeconomic factors like inflation and currency fluctuations on stock performance. The system uses Big Data technologies to handle large datasets, which are imported into MySQL and then transferred to HDFS with Apache Sqoop. Hue provides a graphical interface for easy data analysis without needing complex programming, enabling users to query and visualize stock trends through Hive. The work demonstrates Hadoop's scalability and Hue's user-friendly features for efficient stock market analysis.

Abdelaziz Darwiesh et al. [4] proposes an intelligent risk management system for stock markets, utilizing social media data, particularly Twitter, to enhance risk identification and assessment. The methodology integrates big data analytics and natural language processing (NLP) to analyse user behaviour and sentiments, focusing on risks associated with markets, sessions, and prices. Tools like Twint for data scraping, Spark SQL, and Hadoop for data storage, and advanced NLP techniques are used. The approach identifies risks as financial, operational, or geopolitical, achieving an overall risk analysis accuracy of 73.33%. They have considered the NASDAQ stock market data.

## 4 Methodology

### 4.1 MapReduce

The goal of this MapReduce implementation is to process stock market data for various companies, represented in the form of daily records, and generate aggregated statistics for monthly, and yearly timeframes. The output includes averages and maximum values for key metrics such as close price, volume, open price, high price, and low price.

#### Driver

The Driver is responsible for configuring and executing the MapReduce job.

#### Key Functions

Define the Mapper and Reducer classes, set input and output paths and configure the output key and value types.

#### MapReduce Workflow

##### Mapper

The Mapper processes each row of the dataset and emits key-value pairs based on monthly and yearly groupings.

##### Key Functions

Parse the input line and extract relevant fields (e.g., Ticker, Date, Close, Volume, etc.).

Generate grouping keys:

Monthly Key: MONTHLY:<Ticker> | <YYYY-MM> (e.g., MONTHLY: AAPL | 2023-05).

Yearly Key: YEARLY:<Ticker> | <YYYY> (e.g., YEARLY: AAPL | 2023).

Emit the stock metrics (as a string) for each key.

##### Mapper Output

Key: MONTHLY:<Ticker> | <YYYY-MM> or YEARLY:<Ticker> | <YYYY>.

Value: A CSV-formatted string of metrics: <Close>, <Volume>, <Open>, <High>, <Low>.

##### Reducer

The Reducer aggregates all values for each key emitted by the Mapper. It computes the average and maximum values for Close, Volume, Open, High, and Low.

## Key Functions

Aggregate the metrics:

Sum all values for Close, Volume, Open, High, and Low to calculate averages. Track the maximum values for each metric. Calculate averages by dividing sums by the count of records for the key. Emit the final summarized results.

## Reducer Output

Key: Grouping key (e.g., MONTHLY: AAPL | 2023-05).

Value: Two lines:

Average: Average: Close=<val>, Volume=<val>, Open=<val>, High=<val>, Low=<val>.

Maximum: Max: Close=<val>, Volume=<val>, Open=<val>, High=<val>, Low=<val>.

```
hadoop login: root
root@sandbox.hortonworks.com's password:
root@sandbox.hortonworks.com's password:
Last login: Wed Nov 27 04:13:02 2024 from 172.17.0.3
root@sandbox ~# hadoop fs -get /Project/StockAnalysisDriver.jar
root@sandbox ~# hadoop fs -get /Project/stockdata.csv
get: 'stockdata.csv': file exists
root@sandbox ~# hadoop jar StockAnalysisDriver.jar StockAnalysisDriver /Project/input /Project/output
24/11/27 04:31:02 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
24/11/27 04:31:02 INFO client.AMRProxy: Connecting to ResourceManager at sandbox.hortonworks.com/172.17.0.1:8030
24/11/27 04:31:02 INFO client.AMRProxy: Connecting to Application History server at sandbox.hortonworks.com/172.17.0.2:10200
24/11/27 04:31:02 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and annotate your application with T
oolRunner to remedy this.
24/11/27 04:31:03 INFO Input.FileInputFormat: Total input paths to process : 1
24/11/27 04:31:03 INFO Input.NativeCodeLoader: loaded native gpl library
24/11/27 04:31:03 INFO LocalJobCoordinator: Successfully loaded & initialized native-libs library [hadoop-1.0.0-rcv-7a4b57eece804084320dfb5b7e6a6c8cd4a0]
24/11/27 04:31:04 INFO mapreduce.JobDistributedCache: number of splits:1
24/11/27 04:31:04 INFO mapreduce.JobDistributedCache: Submitting tokens for job: job_1732001003157_0011
24/11/27 04:31:04 INFO impl.YarnClientImpl: Submitted application_1732001003157_0011
24/11/27 04:31:04 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8080/proxy/application_1732001003157_0011/
24/11/27 04:31:04 INFO mapreduce.Job: Running job: job_1732001003157_0011
24/11/27 04:31:04 INFO mapreduce.Job: Job job_1732001003157_0011 running in user mode : false
24/11/27 04:31:04 INFO mapreduce.Job: map 0% reduce 0%
24/11/27 04:31:04 INFO mapreduce.Job: map 100% reduce 0%
24/11/27 04:31:04 INFO mapreduce.Job: map 100% reduce 100%
24/11/27 04:31:04 INFO mapreduce.Job: Job job_1732001003157_0011 completed successfully
24/11/27 04:31:14 INFO mapreduce.Job: Counter: 48
File System Counters
FILE: Number of bytes read:2594424
FILE: Number of bytes written:5477157
FILE: Number of read operations:0
FILE: Number of large read operations:0
FILE: Number of write operations:0
HDFS: Number of bytes read:781205
HDFS: Number of bytes written:141297
HDFS: Number of read operations:0
HDFS: Number of large read operations:0
HDFS: Number of write operations:2
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read:781205
File Output Format Counters
Bytes Written:141297
root@sandbox ~# hadoop fs -cat /Project/output/part-r-00000
MONTHLY:AAPL|2018-05 Average: Close=46.72, Volume=100074000.00, Open=46.81, High=47.00, Low=46.54
Max: Close=46.72, Volume=100074000.00, Open=46.81, High=47.00, Low=46.54
MONTHLY:AAPL|2018-06 Average: Close=47.14, Volume=10010091.43, Open=47.18, High=47.46, Low=46.88
Max: Close=47.14, Volume=10010091.43, Open=47.18, High=47.46, Low=46.88
MONTHLY:AAPL|2018-07 Average: Close=47.54, Volume=74475337.34, Open=47.59, High=47.99, Low=47.29
Max: Close=47.54, Volume=74475337.34, Open=47.59, High=47.99, Low=47.29
MONTHLY:AAPL|2018-08 Average: Close=53.34, Volume=121141269.57, Open=53.68, High=53.64, Low=51.83
Max: Close=53.34, Volume=121141269.57, Open=53.68, High=53.64, Low=51.83
MONTHLY:AAPL|2018-09 Average: Close=53.52, Volume=141666128.32, Open=53.58, High=53.13, Low=53.02
Max: Close=53.52, Volume=141666128.32, Open=53.58, High=53.13, Low=53.02
MONTHLY:AAPL|2018-10 Average: Close=55.21, Volume=155662128.70, Open=55.30, High=55.99, Low=54.47
Max: Close=55.21, Volume=155662128.70, Open=55.30, High=55.99, Low=54.47
MONTHLY:AAPL|2018-11 Average: Close=47.81, Volume=181372765.71, Open=47.96, High=48.40, Low=47.18
Max: Close=47.81, Volume=181372765.71, Open=47.96, High=48.40, Low=47.18
MONTHLY:AAPL|2018-12 Average: Close=41.07, Volume=18720921.16, Open=41.13, High=41.53, Low=40.47
Max: Close=41.07, Volume=18720921.16, Open=41.13, High=41.53, Low=40.47
MONTHLY:AAPL|2019-01 Average: Close=38.54, Volume=15829121.41, Open=38.40, High=38.88, Low=37.90
Max: Close=38.54, Volume=15829121.41, Open=38.40, High=38.88, Low=37.90
MONTHLY:AAPL|2019-02 Average: Close=42.91, Volume=80074922.11, Open=42.85, High=43.21, Low=42.63
Max: Close=42.91, Volume=80074922.11, Open=42.85, High=43.21, Low=42.63
MONTHLY:AAPL|2019-03 Average: Close=43.82, Volume=123462424.78, Open=43.81, High=44.23, Low=43.30
Max: Close=43.82, Volume=123462424.78, Open=43.81, High=44.23, Low=43.30
MONTHLY:AAPL|2019-04 Average: Close=48.11, Volume=95415073.33, Open=48.07, High=48.90, Low=47.64
Max: Close=48.11, Volume=95415073.33, Open=48.07, High=48.90, Low=47.64
MONTHLY:AAPL|2019-05 Average: Close=47.82, Volume=13446608.91, Open=47.74, High=48.34, Low=47.30
Max: Close=47.82, Volume=13446608.91, Open=47.74, High=48.34, Low=47.30
MONTHLY:AAPL|2019-06 Average: Close=52.94, Volume=207001320.80, Open=52.72, High=53.81, Low=52.16
Max: Close=52.94, Volume=207001320.80, Open=52.72, High=53.81, Low=52.16
```

```
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read:781205
File Output Format Counters
Bytes Written:141297
root@sandbox ~# hadoop fs -cat /Project/output/part-r-00000
MONTHLY:AAPL|2018-05 Average: Close=46.72, Volume=100074000.00, Open=46.81, High=47.00, Low=46.54
Max: Close=46.72, Volume=100074000.00, Open=46.81, High=47.00, Low=46.54
MONTHLY:AAPL|2018-06 Average: Close=47.14, Volume=10010091.43, Open=47.18, High=47.46, Low=46.88
Max: Close=47.14, Volume=10010091.43, Open=47.18, High=47.46, Low=46.88
MONTHLY:AAPL|2018-07 Average: Close=47.54, Volume=74475337.34, Open=47.59, High=47.99, Low=47.29
Max: Close=47.54, Volume=74475337.34, Open=47.59, High=47.99, Low=47.29
MONTHLY:AAPL|2018-08 Average: Close=53.34, Volume=121141269.57, Open=53.68, High=53.64, Low=51.83
Max: Close=53.34, Volume=121141269.57, Open=53.68, High=53.64, Low=51.83
MONTHLY:AAPL|2018-09 Average: Close=53.52, Volume=141666128.32, Open=53.58, High=53.13, Low=53.02
Max: Close=53.52, Volume=141666128.32, Open=53.58, High=53.13, Low=53.02
MONTHLY:AAPL|2018-10 Average: Close=55.21, Volume=155662128.70, Open=55.30, High=55.99, Low=54.47
Max: Close=55.21, Volume=155662128.70, Open=55.30, High=55.99, Low=54.47
MONTHLY:AAPL|2018-11 Average: Close=47.81, Volume=181372765.71, Open=47.96, High=48.40, Low=47.18
Max: Close=47.81, Volume=181372765.71, Open=47.96, High=48.40, Low=47.18
MONTHLY:AAPL|2018-12 Average: Close=41.07, Volume=18720921.16, Open=41.13, High=41.53, Low=40.47
Max: Close=41.07, Volume=18720921.16, Open=41.13, High=41.53, Low=40.47
MONTHLY:AAPL|2019-01 Average: Close=38.54, Volume=15829121.41, Open=38.40, High=38.88, Low=37.90
Max: Close=38.54, Volume=15829121.41, Open=38.40, High=38.88, Low=37.90
MONTHLY:AAPL|2019-02 Average: Close=42.91, Volume=80074922.11, Open=42.85, High=43.21, Low=42.63
Max: Close=42.91, Volume=80074922.11, Open=42.85, High=43.21, Low=42.63
MONTHLY:AAPL|2019-03 Average: Close=43.82, Volume=123462424.78, Open=43.81, High=44.23, Low=43.30
Max: Close=43.82, Volume=123462424.78, Open=43.81, High=44.23, Low=43.30
MONTHLY:AAPL|2019-04 Average: Close=48.11, Volume=95415073.33, Open=48.07, High=48.90, Low=47.64
Max: Close=48.11, Volume=95415073.33, Open=48.07, High=48.90, Low=47.64
MONTHLY:AAPL|2019-05 Average: Close=47.82, Volume=13446608.91, Open=47.74, High=48.34, Low=47.30
Max: Close=47.82, Volume=13446608.91, Open=47.74, High=48.34, Low=47.30
MONTHLY:AAPL|2019-06 Average: Close=52.94, Volume=207001320.80, Open=52.72, High=53.81, Low=52.16
Max: Close=52.94, Volume=207001320.80, Open=52.72, High=53.81, Low=52.16
```



```

MONTHLY:NVDA|2023-W2 Average: Close=251.87, Volume=8888116.00, Open=249.75, High=254.60, Low=246.03
Max: Close=277.77, Volume=8485460.00, Open=272.29, High=278.34, Low=271.00
MONTHLY:NVDA|2023-W22 Average: Close=271.49, Volume=8913988.47, Open=271.46, High=274.87, Low=267.53
Max: Close=278.09, Volume=8881359.00, Open=275.86, High=281.10, Low=275.57
MONTHLY:NVDA|2023-W23 Average: Close=310.57, Volume=53171904.55, Open=309.63, High=315.73, Low=304.81
Max: Close=401.13, Volume=354391100.00, Open=405.99, High=419.38, Low=399.40
MONTHLY:QQQ|2018-W2 Average: Close=170.07, Volume=2517248.00, Open=170.13, High=173.20, Low=169.63
Max: Close=170.07, Volume=2517248.00, Open=170.13, High=173.20, Low=169.63
MONTHLY:QQQ|2018-W6 Average: Close=174.80, Volume=37185588.67, Open=174.43, High=178.98, Low=173.47
Max: Close=177.09, Volume=77651470.00, Open=177.09, High=177.90, Low=176.85
MONTHLY:QQQ|2018-W7 Average: Close=177.09, Volume=33410301.23, Open=177.40, High=178.54, Low=176.43
Max: Close=182.42, Volume=8053400.00, Open=181.39, High=182.93, Low=180.29
MONTHLY:QQQ|2018-W8 Average: Close=181.65, Volume=28778536.00, Open=181.34, High=183.23, Low=180.47
Max: Close=188.74, Volume=88198130.00, Open=188.41, High=187.82, Low=185.88
MONTHLY:QQQ|2018-W9 Average: Close=183.35, Volume=39755214.21, Open=183.50, High=184.40, Low=182.48
Max: Close=188.05, Volume=88198130.00, Open=186.10, High=186.49, Low=184.97
MONTHLY:QQQ|2018-W16 Average: Close=174.65, Volume=70513315.48, Open=175.55, High=177.69, Low=172.88
Max: Close=180.17, Volume=141294900.00, Open=186.43, High=187.53, Low=185.70
MONTHLY:QQQ|2018-W31 Average: Close=167.10, Volume=54187730.00, Open=168.99, High=168.61, Low=165.28
Max: Close=170.90, Volume=182350100.00, Open=178.79, High=175.58, Low=173.83
MONTHLY:QQQ|2018-W32 Average: Close=156.46, Volume=72821578.95, Open=159.47, High=161.24, Low=156.41
Max: Close=172.33, Volume=355373500.00, Open=173.10, High=173.33, Low=169.53
MONTHLY:QQQ|2018-W1 Average: Close=161.20, Volume=4418360.00, Open=160.81, High=162.00, Low=159.46
Max: Close=168.16, Volume=74689640.00, Open=166.65, High=168.99, Low=166.47
MONTHLY:QQQ|2019-W4 Average: Close=171.15, Volume=2886873.85, Open=170.81, High=173.77, Low=170.32
Max: Close=173.78, Volume=41818279.00, Open=174.30, High=174.86, Low=173.48
MONTHLY:QQQ|2019-W8 Average: Close=176.84, Volume=37000132.30, Open=176.70, High=177.76, Low=175.66
Max: Close=182.57, Volume=70727100.00, Open=181.76, High=182.85, Low=179.28
MONTHLY:QQQ|2019-W9 Average: Close=186.63, Volume=23483640.29, Open=186.43, High=187.06, Low=185.68
Max: Close=191.02, Volume=3388300.00, Open=181.33, High=191.32, Low=189.84
MONTHLY:QQQ|2019-W5 Average: Close=162.43, Volume=40979723.18, Open=162.30, High=163.71, Low=161.33
Max: Close=193.13, Volume=67185770.00, Open=190.76, High=193.32, Low=189.53
MONTHLY:QQQ|2019-W6 Average: Close=163.00, Volume=34435529.00, Open=163.14, High=164.18, Low=162.01
Max: Close=180.45, Volume=75842390.00, Open=189.72, High=189.77, Low=188.15
MONTHLY:QQQ|2019-W7 Average: Close=162.62, Volume=22471878.18, Open=162.36, High=163.12, Low=161.84
Max: Close=180.20, Volume=41627479.00, Open=188.10, High=198.55, Low=198.43
MONTHLY:QQQ|2019-W8 Average: Close=185.81, Volume=57155612.75, Open=186.05, High=187.58, Low=184.26

```

The above output gives the average of all the features of the dataset excluding the Ticker, Dates with respect to each stock, per every month.

```

WEEKLY:AAPL|2023-W22 Average: Close=187.30, Volume=110651359.30, Open=187.19, High=188.49, Low=185.98
Max: Close=182.01, Volume=426884800.00, Open=182.63, High=182.94, Low=179.12
WEEKLY:AMZN|2023-W22 Average: Close=122.07, Volume=83004993.99, Open=122.13, High=123.63, Low=120.48
Max: Close=186.57, Volume=311345600.00, Open=187.20, High=188.65, Low=184.84
WEEKLY:BRK-B|2023-W22 Average: Close=247.85, Volume=4656395.01, Open=247.92, High=249.72, Low=245.87
Max: Close=359.57, Volume=22303500.00, Open=361.39, High=362.10, Low=355.53
WEEKLY:GOOGL|2023-W22 Average: Close=89.61, Volume=34913868.93, Open=89.56, High=90.60, Low=88.57
Max: Close=149.84, Volume=132345540.00, Open=151.25, High=151.55, Low=148.98
WEEKLY:META|2023-W22 Average: Close=219.33, Volume=23946662.46, Open=219.20, High=222.33, Low=216.27
Max: Close=382.18, Volume=232316600.00, Open=381.68, High=384.33, Low=378.81
WEEKLY:MSFT|2023-W22 Average: Close=208.68, Volume=38164146.52, Open=208.61, High=210.82, Low=206.32
Max: Close=343.11, Volume=110945000.00, Open=344.62, High=349.67, Low=342.20
WEEKLY:NVDA|2023-W22 Average: Close=131.66, Volume=46693580.88, Open=131.56, High=134.15, Low=128.92
Max: Close=401.11, Volume=250920160.00, Open=405.95, High=419.38, Low=399.49
WEEKLY:QQQ|2023-W22 Average: Close=265.77, Volume=48004093.54, Open=265.68, High=267.94, Low=263.27
Max: Close=403.99, Volume=199448100.00, Open=405.57, High=408.71, Low=402.58
WEEKLY:SPY|2023-W22 Average: Close=355.27, Volume=85723022.48, Open=355.23, High=357.49, Low=352.76
Max: Close=477.71, Volume=392220700.00, Open=479.22, High=479.98, Low=476.06
WEEKLY:TSLA|2023-W22 Average: Close=144.41, Volume=136459869.00, Open=144.48, High=147.93, Low=140.81
Max: Close=409.97, Volume=914080943.00, Open=411.47, High=414.50, Low=405.67
WEEKLY:TSM|2023-W22 Average: Close=76.65, Volume=9933653.78, Open=76.70, High=77.56, Low=75.78
Max: Close=140.66, Volume=60793170.00, Open=141.61, High=145.00, Low=139.42
WEEKLY:V|2023-W22 Average: Close=192.90, Volume=8269259.63, Open=192.90, High=194.85, Low=190.93
Max: Close=250.93, Volume=38379570.00, Open=250.05, High=252.67, Low=248.22

```

The above output represents the average of all with respect to every stock for the last week.

```

YEARLY:GOOGL|2020      Average: Close=73.95, Volume=39962114.31, Open=73.88, High=74.83, Low=72.98
      Max: Close=91.25, Volume=108357760.00, Open=91.03, High=92.19, Low=90.85
YEARLY:GOOGL|2021      Average: Close=124.23, Volume=30518727.78, Open=124.15, High=125.26, Low=123.03
      Max: Close=149.84, Volume=97881640.00, Open=149.98, High=150.97, Low=148.90
YEARLY:GOOGL|2022      Average: Close=114.76, Volume=34768810.72, Open=114.88, High=116.48, Low=113.20
      Max: Close=148.00, Volume=123199220.00, Open=151.25, High=151.55, Low=145.52
YEARLY:GOOGL|2023      Average: Close=102.05, Volume=37964794.47, Open=101.63, High=103.21, Low=100.57
      Max: Close=125.05, Volume=119455000.00, Open=125.64, High=126.43, Low=122.74
YEARLY:META|2018       Average: Close=167.73, Volume=25439471.62, Open=167.74, High=169.92, Low=165.61
      Max: Close=217.50, Volume=169059900.00, Open=215.72, High=218.62, Low=214.27
YEARLY:META|2019       Average: Close=181.64, Volume=16153162.00, Open=181.57, High=183.45, Low=179.65
      Max: Close=208.10, Volume=77005780.00, Open=208.67, High=208.93, Low=206.59
YEARLY:META|2020       Average: Close=234.55, Volume=22452856.44, Open=234.35, High=237.94, Low=230.79
      Max: Close=303.91, Volume=76343940.00, Open=300.16, High=304.67, Low=293.05
YEARLY:META|2021       Average: Close=321.17, Volume=18868744.48, Open=321.10, High=325.01, Low=317.41
      Max: Close=382.18, Volume=65654040.00, Open=381.68, High=384.33, Low=378.81
YEARLY:META|2022       Average: Close=180.19, Volume=35587379.88, Open=180.24, High=183.84, Low=176.88
      Max: Close=338.54, Volume=232316600.00, Open=339.95, High=343.09, Low=337.19
YEARLY:META|2023       Average: Close=194.56, Volume=28594959.03, Open=193.65, High=197.05, Low=191.57
      Max: Close=264.72, Volume=150475700.00, Open=265.25, High=268.65, Low=261.29
YEARLY:MSFT|2018       Average: Close=106.68, Volume=31684836.08, Open=106.80, High=107.81, Low=105.51
      Max: Close=115.61, Volume=110945000.00, Open=115.42, High=116.18, Low=114.93
YEARLY:MSFT|2019       Average: Close=130.38, Volume=24533254.25, Open=130.34, High=131.23, Low=129.31
      Max: Close=158.96, Volume=55442870.00, Open=159.45, High=159.55, Low=158.22
YEARLY:MSFT|2020       Average: Close=193.03, Volume=37668574.98, Open=192.91, High=195.47, Low=190.38
      Max: Close=231.65, Volume=97073560.00, Open=229.27, High=232.86, Low=227.35
YEARLY:MSFT|2021       Average: Close=275.94, Volume=26014353.21, Open=275.67, High=278.02, Low=273.44
      Max: Close=343.11, Volume=69870640.00, Open=344.62, High=349.67, Low=342.20
YEARLY:MSFT|2022       Average: Close=268.92, Volume=31224383.59, Open=269.11, High=272.50, Low=265.29
      Max: Close=334.75, Volume=90428850.00, Open=335.35, High=338.00, Low=329.78
YEARLY:MSFT|2023       Average: Close=273.96, Volume=30891630.10, Open=273.45, High=276.52, Low=270.87
      Max: Close=332.89, Volume=69527370.00, Open=335.23, High=335.94, Low=330.52
YEARLY:NVDA|2018       Average: Close=57.57, Volume=50047740.43, Open=57.71, High=58.59, Low=56.64
      Max: Close=72.34, Volume=196162800.00, Open=72.33, High=73.19, Low=71.40
YEARLY:NVDA|2019       Average: Close=43.65, Volume=45592350.37, Open=43.61, High=44.25, Low=43.01
      Max: Close=59.84, Volume=250920160.00, Open=60.13, High=60.45, Low=59.60

```

The above output represents the average of all the features of the stock with respect to the year.

## 4.2 Sqoop

Sqoop stands for SQL-to-Hadoop, is an open-source tool in hadoop ecosystem which is designed for efficiently transferring bulk data in between Apache Hadoop and Structured databases which includes MySQL, PostgreSQL, and Oracle. In this project we have utilized the tool Sqoop, to import the data from the Relational Database Management Systems like MySQL into Hadoop Distributed File System and integrate this imported data into the Hive tool for execution of the queries on the data for betterment of analysing the data. In the below we can see the implementation of importing the data into Hive from the Structured database by leveraging the Sqoop tool.

Initially to connect with the mysql in hadoop we need to use the command of - mysql -u root -h sandbox.hortonworks.com -p, this allows to to connect with the mysql and after executing this command it asks for the password and it is, hadoop. After execution you will be entering into the mysql environment. Below you can see the databases and tables present in a particular database.

```
mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| hive |
| mysql |
| performance_schema |
| project |
| ranger |
+-----+
6 rows in set (0.00 sec)
```

```
mysql> use project;
Database changed
mysql> show tables;
+-----+
| Tables_in_project |
+-----+
| stock |
+-----+
1 row in set (0.00 sec)
```

The data has already been in the project named database, of the table names stocks, by entering the command - select \* from stocks limit 5. This command displays a total of 5 rows which are at the top of the table.

```
mysql> LOAD DATA INFILE '/var/lib/mysql-files/stockdata.csv' INTO TABLE stocks FIELDS TERMINATED
BY ',' ENCLOSED BY '"' LINES TERMINATED BY '\n' IGNORE 1 LINES;
Query OK, 15108 rows affected, 628 warnings (1.16 sec)
Records: 15108 Deleted: 0 Skipped: 0 Warnings: 628

mysql> select * from stocks limit 5;
+-----+-----+-----+-----+-----+-----+-----+
| Ticker | Dates      | Close | Volume | Open  | High  | Low   |
+-----+-----+-----+-----+-----+-----+-----+
| AAPL   | 05/31/2023 | 177.25 | 99625290 | 177.33 | 179.35 | 176.76 |
| AAPL   | 05/30/2023 | 177.30 | 55964400 | 176.96 | 178.99 | 176.57 |
| AAPL   | 05/26/2023 | 175.43 | 54834980 | 173.32 | 175.77 | 173.11 |
| AAPL   | 05/25/2023 | 172.99 | 56058260 | 172.41 | 173.90 | 171.69 |
| AAPL   | 05/24/2023 | 171.84 | 45143490 | 171.09 | 172.42 | 170.52 |
+-----+-----+-----+-----+-----+-----+-----+
5 rows in set (0.01 sec)
```

For the implementation of the Sqoop, exit from the mysql environment. Now in the normal root of SSH Client, you need to enter the below commands for the migration



```

[root@sandbox ~]# sqoop import \
> --connect jdbc:mysql://localhost/project \
> --username root \
> --password hadoop \
> --table stocks \
> --driver com.mysql.jdbc.Driver \
> --hive-import \
> --hive-table stockmarket.detail \
> --m 1
Warning: /usr/hdp/2.5.0.0-1245/accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
24/11/29 13:42:58 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6.2.5.0.0-1245
24/11/29 13:42:58 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure.
Consider using -P instead.
24/11/29 13:42:58 INFO tool.BaseSqoopTool: Using Hive-specific delimiters for output. You can override
delimiters with --fields-terminated-by, etc.
24/11/29 13:42:58 WARN sqoop.ConnFactory: Parameter --driver is set to an explicit driver however
appropriate connection manager is not being set (via --connection-manager). Sqoop is going to fall
back to org.apache.sqoop.manager.GenericJdbcManager. Please specify explicitly which connection
manager should be used next time.
24/11/29 13:42:58 INFO manager.SqlManager: Using default fetchSize of 1000
24/11/29 13:42:58 INFO tool.CodeGenTool: Beginning code generation
24/11/29 13:42:59 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM stocks AS t WHERE 1=0
24/11/29 13:42:59 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM stocks AS t WHERE 1=0
24/11/29 13:42:59 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/hdp/2.5.0.0-1245/hadoop-mapreduce
Note: /tmp/sqoop-root/compile/624968129ea2db1436685dc1adbda810/stocks.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
24/11/29 13:43:05 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-root/compile/624968129ea2db1436685dc1adbda810/stocks.jar
24/11/29 13:43:05 INFO mapreduce.ImportJobBase: Beginning import of stocks
24/11/29 13:43:06 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM stocks AS t WHERE 1=0

```

of the data from mysql to Hive. By utilizing the sqoop, we would be connecting to the mysql database by entering our database name at the end, for every sandbox execution environments the username will be root, and the password is hadoop, update the table name, create a hive database and a table with column values in it and give the databasename.tablename as the location to store the imported data, at final declare how many mappers should be executed for this process.

```

File: Number of bytes read=0
FILE: Number of bytes written=162691
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=87
HDFS: Number of bytes written=779194
HDFS: Number of read operations=4
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Other local map tasks=1
  Total time spent by all maps in occupied slots (ms)=12972
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=12972
  Total vcore-milliseconds taken by all map tasks=12972
  Total megabyte-milliseconds taken by all map tasks=3243000
Map-Reduce Framework
  Map input records=15108
  Map output records=15108
  Input split bytes=87
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=193
  CPU time spent (ms)=2360
  Physical memory (bytes) snapshot=145682432
  Virtual memory (bytes) snapshot=1934544896
  Total committed heap usage (bytes)=35651584
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=779194
24/11/29 13:43:59 INFO mapreduce.ImportJobBase: Transferred 760.9316 KB in 52.6457 seconds (14.45
KB/sec)
24/11/29 13:43:59 INFO mapreduce.ImportJobBase: Retrieved 15108 records.

```

```

24/11/29 13:47:03 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM stocks AS t W
HERE 1=0
24/11/29 13:47:03 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM stocks AS t W
HERE 1=0
24/11/29 13:47:03 WARN hive.TableDefWriter: Column Close had to be cast to a less precise type in
Hive
24/11/29 13:47:03 WARN hive.TableDefWriter: Column Open had to be cast to a less precise type in
Hive
24/11/29 13:47:03 WARN hive.TableDefWriter: Column High had to be cast to a less precise type in
Hive
24/11/29 13:47:03 WARN hive.TableDefWriter: Column Low had to be cast to a less precise type in H
ive
24/11/29 13:47:03 INFO hive.HiveImport: Loading uploaded data into Hive

Logging initialized using configuration in jar:file:/usr/hdp/2.5.0.0-1245/hive/lib/hive-common-1.
2.1000.2.5.0.0-1245.jar!/hive-log4j.properties
OK
Time taken: 2.893 seconds
Loading data to table stockmarket.detail
Table stockmarket.detail stats: [numFiles=1, numRows=0, totalSize=779194, rawDataSize=0]
OK
Time taken: 2.844 seconds

```

```

hive> show databases;
OK
default
foodmart
project
stock
stockmarket
xademo
Time taken: 1.596 seconds, Fetched: 6 row(s)
hive> use stockmarket;
OK
Time taken: 0.323 seconds
hive> select * from detail limit 5;
OK
AAPL    05/31/2023    177.25  99625290    177.33  179.35  176.76
AAPL    05/30/2023    177.3   55964400    176.96  178.99  176.57
AAPL    05/26/2023    175.43  54834980    173.32  175.77  173.11
AAPL    05/25/2023    172.99  56058260    172.41  173.9   171.69
AAPL    05/24/2023    171.84  45143490    171.09  172.42  170.52
Time taken: 0.598 seconds, Fetched: 5 row(s)

```

In case if the commands or the names of the database or tables or the locations are given incorrectly then the execution of importing the data via sqoop might fail. In case if all the details have been perfectly matched then the sqoop operation takes place. As you can see in the above ones, it was present in the hive environment and the data has been successfully loaded into the table named detail and of the database named project in the hive environment. Now utilizing this data we would be performing some analysis in the Hive environment.

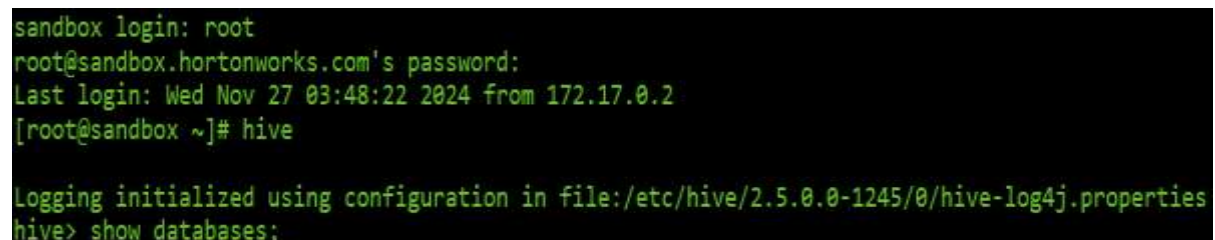
### 4.3 Hive

Apache Hive is a data warehouse which is built on top of Hadoop. It is utilized for managing and querying large amounts of datasets in distributed storage systems. Hive helps in providing an abstraction which utilizes MapReduce, which simplifies the process of querying large amounts of datasets. Apache Hive offers a high-level interface for the users which enables them to write the queries like the queries we execute in SQL. Hive was developed by Facebook.

In this project we have utilized Hive as a central data warehouse system which helps in managing the data, querying and processing large datasets. The dataset we included has a total of 12 different types of stocks. The features are Ticker - describes the name of the stock, Dates - describes the timestamp with respect to the stock, Open - describes the price of the stock at the opening session of the market, Close - describes the price of the stock at the closing session of the market, High - describes the highest price of the stock with respect to each stock every day, Low - describes the lowest price of the stock during the trading session of each day and Volume - describes about the total number of shares traded with respect to the stock per day.

Since Hive provides a distributed storage and processing system, it will be effective for handling large volumes of structured data, in the perspective of our project, we have included the Stocks dataset into the Hadoop environment. The dataset consists of around 15000 records of the data. It consists of 5 years of data from 2018 to 2023. Hive scalability helps in good management and processing of this data efficiently. By utilizing the CLI - Command Line Interface provided by the Hive environment which will be like the SQL-like query language, makes it easier for analysing large datasets equipped with MapReduce.

We have included the entire process of the Hive operations we have performed in the Hadoop environment. We have accessed the tools by using the Sandbox Hortonworks platforms. We have loaded the dataset into the HDFS. Open the SSH Client and login into your account by entering the details. Later, type the hive in the command interface so that the environment changes from the normal to hive interface.



```
sandbox login: root
root@sandbox.hortonworks.com's password:
Last login: Wed Nov 27 03:48:22 2024 from 172.17.0.2
[root@sandbox ~]# hive

Logging initialized using configuration in file:/etc/hive/2.5.0.0-1245/0/hive-log4j.properties
hive> show databases;
```

After clicking the hive command, we can see that the interface has been changed from the [root@sandbox ~] # to hive>. This Hive acts so like the SQL interface. First, we need to know the databases present in the environment, if we require any database with a new one, we can create a new database. For knowing the existing databases, we have the command hive> show databases; this command displays all the existing databases present in the Hive instance. For creating a new database, we can use the



command `hive> create 'database database_name'`. This creates a new database in the Hive instance. I have created a new database without merging all the files into one database.

```
hive> show databases;
OK
default
foodmart
project
stock
stockmarket
xademo
Time taken: 2.123 seconds, Fetched: 6 row(s)
hive> use stockmarket;
OK
Time taken: 0.54 seconds
```

I have created a database with name stock market. For performing operations in the database, we use the command `hive> 'use database_name'`. This command allows us to use the database and perform operations within the database. For knowing the existing tables residing in the database we use command `hive> 'show tables'`. This command lists out all the tables present in the database. Since its a newly created database there aren't any tables in the database. For creating a table, we use the command `hive> 'create table table_name(column1_name datatype, column2_name datatype, column3_name datatype, .....);'`. Here we are loading a stocks dataset into the table, which is in csv format, so we add additional lines for the above query with including 'row format delimited fields terminated by',' and 'TBLPROPERTIES ('skip.header.line.count'='1');

```
hive> create table details(Ticker STRING, Dates STRING,Close DECIMAL,Volume INT,Open DECIMAL,High DECIMAL,Low DECIMAL)
> ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
> TBLPROPERTIES ("skip.header.line.count"="1");
OK
Time taken: 0.97 seconds
```

The above lines indicate – row format refers to how individual records (or rows) of data are organized or stored. Delimited fields refers that the data in each row is broken into fields (or columns) based on a specific delimiter. A delimiter is a character that marks the boundary between different fields (or data values) within the row. Terminated by ',': This part specifies that the delimiter separating the fields in a row is a comma (,). Upon creation of the table to check the column fields and datatypes are correctly specified, we can use the command `hive> 'describe table_name'`. This displays the columns and the related datatypes of the column which are present in the table.

```
hive> describe details;
OK
ticker                string
dates                 string
close                 decimal(10,0)
volume                int
open                  decimal(10,0)
high                  decimal(10,0)
low                   decimal(10,0)
Time taken: 0.841 seconds, Fetched: 7 row(s)
```

After that by using the command `hive> load data local inpath 'dataset path' into table table_name`; we can load the dataset file into the table which is created in the database of Hive instance. For every successful execution of the command, we can check the OK message which is being projected on the screen.

```
hive> LOAD DATA LOCAL INPATH 'stockdata.csv' INTO TABLE details;
Loading data to table stockmarket.details
Table stockmarket.details stats: [numFiles=1, numRows=8, totalSize=783566, rawDataSize=8]
OK
Time taken: 1.924 seconds
hive> SELECT * FROM details LIMIT 5;
OK
AAPL  05/31/2023    177    99625290    177    179    177
AAPL  05/30/2023    177    55864400    177    179    177
AAPL  05/26/2023    175    54834980    173    176    173
AAPL  05/25/2023    173    56058260    172    174    172
AAPL  05/24/2023    172    45143490    171    172    171
Time taken: 0.385 seconds, Fetched: 5 row(s)
```

We can check if the data is correctly loaded or not by using the command `hive> 'SELECT * FROM table_name'`, this command lists out all the records which are present in the table. In case if you need to display a few numbers of records of the table we can use the command, `hive> 'SELECT * FROM table_name limit 10(number of records)'`. This command displays only the top 10 records of the table. In the above I have passed the query to display the top 5 records of the table. Up to now this is the process of the changing environment to hive, creating and managing the database, table creation and loading of the dataset into tables. For making some decisions on the considered data we need to perform some analysis which helps in understanding the existing data better.

For processing the table contents, we first checked the total number of records associated with each unique stock name. So, we have used the below command from the image to return the total number of records associated with each stock name in the table. As you can see the query returns the output of the number of records associated with each stock name. There is a total of 12 stocks present in the table and each stock has the same number of records associated with value of 1259.



```

hive> select ticker, count(*) as record_count
> from details
> group by ticker;
Query ID = root_20241127043414_0807fdd5-74d8-4c6b-b53c-86c99c2311d1
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.

Status: Running (Executing on YARN cluster with App id application_1732601083157_0015)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED    1         1         0         0         0         0
Reducer 2 ..... SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 25.98 s
-----
OK
AAPL      1259
AMZN      1259
BRK-B     1259
GOOGL     1259
META      1259
MSFT      1259
NVDA      1259
QQQ       1259
SPY       1259
TSLA      1259
TSM       1259
V         1259
Time taken: 44.364 seconds, Fetched: 12 row(s)
hive>

```

For further we have tried to analyse the Closing value of each stock in the table. We found the highest and lowest closing prices for each of the stock. This query groups the data by each stock (ticker symbol) and calculates the maximum and minimum closing prices (close) for that stock. This helps identify the price range a stock has traded within. Such insights are useful for investors or analysts to understand stock performance over time.

```

hive> select ticker, MAX(Close) as max_close_price, MIN(Close) as min_close_price
> from details
> group by ticker;
Query ID = root_20241127043634_649470a8-10b9-4cf3-8116-61d867bf4268
Total jobs = 1
Launching Job 1 out of 1

Status: Running (Executing on YARN cluster with App id application_1732601083157_0015)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED    1         1         0         0         0         0
Reducer 2 .....  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100%  ELAPSED TIME: 11.99 s
-----
OK
AAPL    182    36
AMZN    187    67
BRK-B   360   162
GOOGL   150    49
META    382    89
MSFT    343    94
NVDA    401    32
QQQ     404   144
SPY     478   223
TSLA    410    12
TSM     141    34
V       251   122
Time taken: 13.578 seconds, Fetched: 12 row(s)

hive> select year,ticker,avg_price_range
> from(select YEAR(FROM_UNIXTIME(UNIX_TIMESTAMP(dates,'MM/dd/yyyy')))AS year,
> ticker,
> AVG(high-low) AS avg_price_range,
> ROW_NUMBER() OVER (PARTITION BY YEAR(FROM_UNIXTIME(UNIX_TIMESTAMP(dates,'MM/dd/yyyy'))) ORDER BY AVG(high-low) DESC) AS rank
> FROM details
> GROUP BY YEAR(FROM_UNIXTIME(UNIX_TIMESTAMP(dates,'MM/dd/yyyy'))),ticker
> ) ranked_data
> WHERE rank=1;
Query ID = root_20241127050106_f778d3c4-f9a4-4de8-81c5-c8a1e90a2e7b
Total jobs = 1
Launching Job 1 out of 1
Ter session was closed. Reopening...
Session re-established.

Status: Running (Executing on YARN cluster with App id application_1732601083157_0016)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED    1         1         0         0         0         0
Reducer 2 .....  SUCCEEDED    1         1         0         0         0         0
Reducer 3 .....  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 03/03 [=====>>>] 100%  ELAPSED TIME: 37.74 s
-----
OK
NULL    TSLA    7.4435
2018    META    4.4783
2019    META    3.8212
2020    META    7.0129
2021    TSLA    10.8289
2022    TSLA    13.7815
2023    NVDA    9.1774
Time taken: 67.327 seconds, Fetched: 7 row(s)
hive>

```

The above query identifies the stocks name column ticker with the highest average daily price range (high - low) for each year in the dataset. It calculates the average price range for each ticker by grouping the data by year and ticker and using the AVG () function. The ROW\_NUMBER () function ranks tickers within each year based on their average price range, with rank 1 assigned to the highest. The outer query filters the results to only include the top-ranked ticker (highest average price range) for each

year. This query helps analysts determine the most volatile stock annually, providing insights into price movement trends and potential opportunities for traders.

```
hive> select d.ticker,d.dates,d.open from details d
> JOIN(SELECT ticker, MAX(open) AS max_open
> FROM details
> GROUP BY ticker
> ) m
> ON d.ticker=m.ticker AND d.open=m.max_open;
Query ID = root_20241127052401_8a865287-3d89-4c64-9b8e-197a83ff0d92
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
```

```
Status: Running (Executing on YARN cluster with App id application_1732601083157_0017)
```

	VERTICES	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....		SUCCEEDED	1	1	0	0	0	0
Map 3 .....		SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....		SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 29.06 s
```

```
OK
```

AAPL	01-04-2022	183
AMZN	07-12-2021	187
BRK-B	03/29/2022	361
BRK-B	03/28/2022	361
GOOGL	02-02-2022	151
META	09/13/2021	382
META	09-02-2021	382
MSFT	11/22/2021	345
NVDA	05/30/2023	406
QQQ	11/22/2021	406
SPY	01-04-2022	479
TSLA	11-04-2021	411
TSM	02/16/2021	142
V	07/28/2021	250
V	07/16/2021	250

The above query retrieves the ticker, date, and opening price for the record with the highest opening price (Open - column name) for each stock ticker in the dataset. By using a subquery to calculate the maximum opening price for each ticker, and then joining it back with the original dataset, it ensures that the full record corresponding to this maximum value is fetched. This is particularly useful for identifying the day each stock opened at its highest price, which can provide valuable insights for historical analysis, peak performance evaluation, and identifying trading opportunities.

```

FAILED: SemanticException [Error 10004]: Line 1:252 Invalid table alias or column reference
hive> select d.ticker,d.dates,d.Close from details d
> JOIN(SELECT ticker,MAX(Close) AS max_close
> FROM details
> GROUP BY ticker) m
> ON d.ticker=m.ticker AND d.open=m.max_close;
Query ID = root_20241129151217_d2787099-65dd-473b-8112-312728596c16
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.

Status: Running (Executing on YARN cluster with App id application_1732883356534_0008)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED   1         1         0         0         0         0
Map 3 .....  SUCCEEDED   1         1         0         0         0         0
Reducer 2 ..... SUCCEEDED   1         1         0         0         0         0
-----
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 34.44 s
-----
OK
AMZN    07-12-2021    186
GOOGL   11/19/2021    149
GOOGL   11-08-2021    149
META    09/13/2021    377
META    09-02-2021    375
MSFT    12/28/2021    341
MSFT    11/19/2021    343
QQQ     12/28/2021    402
SPY     12/30/2021    476
SPY     12/28/2021    477
TSM     01/13/2022    139
Time taken: 43.892 seconds, Fetched: 11 row(s)

```

The above query returns the highest value of the closing price with respect to every stock included in the dataset, displaying the timestamp that on which date, that the stock closing price is high.

```

hive> SELECT d.ticker,d.dates,d.high, YEAR(FROM_UNIXTIME(UNIX_TIMESTAMP(d.dates,'MM/dd/yyyy')) AS year
> FROM details d
> JOIN (
> SELECT
> ticker,
> YEAR(FROM_UNIXTIME(UNIX_TIMESTAMP(dates,'MM/dd/yyyy')) AS year,
> MAX(high) AS max_high
> FROM details
> GROUP BY ticker, YEAR(FROM_UNIXTIME(UNIX_TIMESTAMP(dates,'MM/dd/yyyy'))
> ) m
> ON d.ticker=m.ticker
> AND YEAR(FROM_UNIXTIME(UNIX_TIMESTAMP(d.dates,'MM/dd/yyyy'))=m.year
> AND d.high=m.max_high;
Query ID = root_20241127053628_33cb3aff-0bb3-488b-b89a-ce27d42a372a
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.

Status: Running (Executing on YARN cluster with App id application_1732601083157_0018)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED   1         1         0         0         0         0
Map 3 .....  SUCCEEDED   1         1         0         0         0         0
Reducer 2 ..... SUCCEEDED   1         1         0         0         0         0
-----
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 23.57 s
-----
OK

```



AMZN	09/28/2018	101	2018
AMZN	09/27/2018	101	2018
AMZN	08/31/2018	101	2018
AMZN	08/30/2018	101	2018
AMZN	07/17/2019	101	2019
AMZN	07/16/2019	101	2019
AMZN	07/15/2019	101	2019
AMZN	10/13/2020	175	2020
AMZN	08/31/2020	175	2020
AMZN	07/13/2021	189	2021
AMZN	03/29/2022	171	2022
AMZN	05/30/2023	123	2023
BRK-B	09/21/2018	223	2018
BRK-B	09/20/2018	223	2018
BRK-B	12/20/2019	228	2019
BRK-B	12/17/2019	228	2019
BRK-B	12/16/2019	228	2019
BRK-B	11/24/2020	235	2020
BRK-B	11/17/2020	235	2020
BRK-B	12/30/2021	302	2021
BRK-B	12/16/2021	302	2021

The above query helps in identifying the highest price (high) recorded by each stock ticker for every year in the dataset, along with the date it occurred. It achieves this by first calculating the maximum high price for each ticker and year combination using a subquery. The main query then joins this result back with the original dataset to retrieve the full record (including the date and price) corresponding to these maximum values. This query is particularly useful for annual performance analysis of stocks, allowing analysts or investors to pinpoint peak price movements for each stock in any given year. It provides valuable insights into market trends, helps identify the most profitable trading days, and offers context for historical comparisons across years. The results can also be used to assess volatility or the impact of specific events on stock performance in a particular year.

#### 4.4 Pyspark

##### Importing packages:

In python, to use the spark module we need to install the pyspark package using the pip command.

PySpark is a Python API that allows users to interact with Apache Spark, an open-source distributed computing framework

pip install pyspark

#### **4.4.1 Loading of Data**

##### **Creating spark session**

SparkSession is an entry point into all functionality in Spark, and it is required if we want to build dataframe operations, SQL queries, and machine learning tasks in PySpark.

appName specifies a name for the application, in this case, "Stock Price Analysis." This helps in identifying the application when monitoring tasks in the Spark UI.

getOrCreate ensures that a Spark session is either created if it doesn't already exist or retrieves an existing one.

##### **Reading the Stock Data**

read.csv () Specifies the path to the directory or file containing the stock data.

show () displays the first five rows of the DataFrame in a tabular format, helping to visually inspect the data.

printSchema() defines the structure and data types of each column.

#### **4.4.2 Data Preprocessing**

##### **Parsing Date Column**

To ensure accurate analysis, the stock data underwent a series of data cleaning steps using PySpark. These steps involved parsing dates, converting numerical columns, and changing data types for consistency.

The ParsedDate column is created by converting the string format ("MM/DD/YYYY") into a DateType using a user-defined function (UDF). This allows for efficient date-based operations such as filtering, aggregation, and sorting.

##### **2.3.2 Parsing Numerical Columns**

Many numerical columns in the dataset, such as Open, Close/Last, Low, and High, contained string representations (e.g., "\$123.45") or required normalization. These were cleaned and converted to FloatType

The num\_parser function strips dollar signs and converts the values into floats.

The withColumn() method was used to apply this transformation to each relevant column.

This ensures that the columns are numeric, facilitating mathematical operations like calculations and aggregations.

##### **Converting Volume to Integer Type**

The Volume column, often stored as a string, needs to be converted to an IntegerType to align with its numerical nature.

A UDF (parse\_int) is created to parse and convert values to integers.

The .withColumn() method is used to apply this transformation.

After the data cleaning procedures, the schema was inspected to confirm the successful transformations:

Dates were parsed and stored as DateType.

Numerical columns like Open, Close, Low, High, and Volume were appropriately typed as FloatType or IntegerType.

The data cleaning steps ensures

- Proper data types for analysis.
- Improved data consistency and usability.
- Elimination of formatting issues, enabling seamless integration into analytical workflows.

#### 4.4.3 Exploratory Data Analysis

After cleaning the stock data, EDA techniques were applied to understand the dataset's characteristics and derive meaningful insights.

summary	Volume	Open	Low	High	Close
count	15108	15108	15108	15108	15108
mean	5.1868408793685466E7	180.09656566181036	177.9982781513109	182.1253348687101	180.1256089860054
stddev	5.496484129953464E7	101.16125813324383	100.26590135955234	101.96625521621753	101.14891782168543
min	961133	12.07	11.8	12.45	11.93
max	914080943	479.22	476.06	479.98	477.71

The method describe () is used to provide the basic descriptive statistics such as count, mean, standard deviation, minimum, and maximum values for selected columns (Volume, Open, Low, High, and Close).

This helps in understanding the data's central tendencies, spread, and outliers.

#### Calculating Maximum Stock Price for Various Stocks

```
cleaned_stocks.groupBy("Ticker").max("Open"). show (15)
```

+-----+-----+	
Ticker	max(Open)
+-----+-----+	
BRK-B	361.39
AAPL	182.63
META	381.68
TSLA	411.47
AMZN	187.2
MSFT	344.62
TSM	141.61
QQQ	405.57
V	250.05
GOOGL	151.25
SPY	479.22
NVDA	405.95
+-----+-----+	

Groups the data by the Ticker column (stock identifier) and calculates the maximum opening price.

This highlights the peak stock price for each stock, useful for identifying top-performing stocks.

### **Aggregating Stock Statistics**



```
cleaned_stocks.groupBy("Ticker").agg(
  func.max("Open").alias("MaxStockPrice"),
  func.sum("Volume").alias("TotalVolume")
).show(15)
```

Ticker	MaxStockPrice	TotalVolume
BRK-B	361.39	5862401321
AAPL	182.63	139310061360
META	381.68	30148848043
TSLA	411.47	171802975076
AMZN	187.2	104503287430
MSFT	344.62	37976660472
TSM	141.61	12506470104
QQQ	405.57	60437153773
V	250.05	10410997871
GOOGL	151.25	43956560981
SPY	479.22	107925285300
NVDA	405.95	58787218324

`func.max("Open")`: Computes the maximum opening price (`MaxStockPrice`) for each stock.

`func.sum("Volume")`: Aggregates the total trading volume (`TotalVolume`) for each stock.

Combines price and volume metrics, providing a more comprehensive view of stock performance.

### Date Manipulation for Yearly Maximum Prices

```
cleaned_stocks=(cleaned_stocks.withColumn("Year",
func.year(cleaned_stocks.ParsedDate)).withColumn("Month",func.month(cleaned_st
ocks.ParsedDate)).withColumn("Day",func.dayofmonth(cleaned_stocks.ParsedDate)
).withColumn("Week", func.weekofyear(cleaned_stocks.ParsedDate)))
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| Ticker | ParsedDate | Volume | Open | Low | High | Close | Year | Month | Day | Week |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| BRK-B | 2023-05-31 | 6175417 | 321.12 | 319.39 | 322.41 | 321.08 | 2023 | 5 | 31 | 22 |
| BRK-B | 2023-05-30 | 3232461 | 321.86 | 319.0 | 322.47 | 322.19 | 2023 | 5 | 30 | 22 |
| BRK-B | 2023-05-26 | 3229873 | 320.44 | 319.67 | 322.63 | 320.6 | 2023 | 5 | 26 | 21 |
| BRK-B | 2023-05-25 | 4251935 | 320.56 | 317.71 | 320.56 | 319.02 | 2023 | 5 | 25 | 21 |
| BRK-B | 2023-05-24 | 3075393 | 322.71 | 319.56 | 323.0 | 320.2 | 2023 | 5 | 24 | 21 |
| BRK-B | 2023-05-23 | 4031342 | 328.19 | 322.97 | 329.27 | 323.11 | 2023 | 5 | 23 | 21 |
| BRK-B | 2023-05-22 | 2763422 | 330.75 | 328.35 | 331.49 | 329.13 | 2023 | 5 | 22 | 21 |
| BRK-B | 2023-05-19 | 4323538 | 331.0 | 329.12 | 333.94 | 330.39 | 2023 | 5 | 19 | 20 |
| BRK-B | 2023-05-18 | 2808329 | 326.87 | 325.85 | 329.98 | 329.76 | 2023 | 5 | 18 | 20 |
| BRK-B | 2023-05-17 | 3047626 | 325.02 | 324.82 | 328.26 | 327.39 | 2023 | 5 | 17 | 20 |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 10 rows

```

Year: Extracts the year from the ParsedDate column.

Month: Extracts the month for monthly analysis.

Day: Extracts the specific day of the month.

Week: Identifies the week number of the year.

This enables time-based analysis, such as finding the maximum stock price for each year or observing seasonal trends.

## Yearly Stock Price Analysis

```

yearly = cleaned_stocks.groupby(['Ticker', 'Year']).agg(func.max("Open").alias("YearlyHigh"), func.min("Open").alias("YearlyLow"))
yearly.show()

```

```

+-----+-----+-----+-----+
| Ticker | Year | YearlyHigh | YearlyLow |
+-----+-----+-----+-----+
| META | 2020 | 300.16 | 139.75 |
| BRK-B | 2023 | 331.0 | 294.68 |
| MSFT | 2019 | 159.45 | 99.55 |
| MSFT | 2021 | 344.62 | 212.17 |
| BRK-B | 2018 | 224.0 | 185.43 |
| META | 2021 | 381.68 | 247.9 |
| TSLA | 2019 | 29.0 | 12.07 |
| META | 2018 | 215.72 | 123.1 |
| AAPL | 2020 | 138.05 | 57.02 |
| MSFT | 2020 | 229.27 | 137.01 |
| TSLA | 2021 | 411.47 | 184.18 |
| BRK-B | 2021 | 300.88 | 228.21 |
| TSLA | 2018 | 25.0 | 17.02 |
| MSFT | 2018 | 115.42 | 95.14 |
| AMZN | 2020 | 177.35 | 82.08 |
| META | 2022 | 339.95 | 90.08 |
| AAPL | 2022 | 182.63 | 127.99 |
| META | 2019 | 208.67 | 128.99 |
| TSLA | 2020 | 233.33 | 24.98 |
| AMZN | 2022 | 170.44 | 82.8 |
+-----+-----+-----+-----+
only showing top 20 rows

```

`func.max("Open")`: Calculates the highest opening price for each stock in each year (YearlyHigh).

`func.min("Open")`: Calculates the lowest opening price for each stock in each year (YearlyLow).

A summary of the annual performance of each stock. This highlights the range of stock price fluctuations over the year, helping identify high-performing and volatile stocks.

### Monthly Stock Price Analysis

```
monthly = cleaned_stocks.groupby(['Ticker', 'Year', 'Month']).agg(  
    func.max("Open").alias("MonthHigh"),  
    func.min("Open").alias("MonthLow")  
)  
monthly.show()
```

Ticker	Year	Month	MonthHigh	MonthLow
BRK-B	2022	10	297.98	260.58
META	2020	6	241.28	209.75
BRK-B	2018	9	222.13	209.21
MSFT	2022	6	275.2	243.86
MSFT	2021	2	245.03	230.01
MSFT	2020	1	174.05	157.08
BRK-B	2021	10	290.85	273.02
BRK-B	2020	10	216.74	200.03
TSLA	2023	4	199.91	152.64
TSLA	2019	4	19.22	15.72
AMZN	2018	5	81.15	81.15
MSFT	2019	6	137.45	121.28
META	2023	4	239.89	208.84
AMZN	2020	2	108.65	90.73
MSFT	2022	10	247.93	219.85
AMZN	2022	3	170.38	136.68
TSLA	2020	10	151.48	135.63
BRK-B	2019	9	212.24	201.19
TSLA	2021	12	386.9	303.57
BRK-B	2021	6	292.91	275.0

only showing top 20 rows

`func.max("Open")`: Calculates the highest opening price for each stock in each month (MonthHigh).

`func.min("Open")`: Calculates the lowest opening price for each stock in each month (MonthLow).

This captures monthly trends and seasonality in stock prices and helps identify periods of significant price changes or seasonal patterns.

### Weekly Stock Price Analysis

```
weekly = cleaned_stocks.groupby(['Ticker', 'Year', 'Week']).agg(  
    func.max("Open").alias("WeekHigh"),  
    func.min("Open").alias("WeekLow")  
)  
weekly.show()
```

Ticker	Year	Week	WeekHigh	WeekLow
BRK-B	2022	14	352.0	341.17
BRK-B	2022	10	326.59	322.49
BRK-B	2021	14	264.22	260.02
META	2022	43	131.68	97.98
META	2020	6	212.51	203.44
TSLA	2022	20	255.72	235.67
TSLA	2020	19	52.92	46.73
TSLA	2020	16	51.49	39.34
TSLA	2018	39	20.86	18.02
BRK-B	2018	48	217.23	209.3
MSFT	2022	6	309.87	301.25
MSFT	2021	2	218.47	213.52
META	2022	40	140.49	136.76
AAPL	2020	27	91.96	88.31
BRK-B	2020	19	180.05	173.4
MSFT	2020	1	158.78	158.32
META	2020	36	298.88	287.25
AAPL	2021	25	134.45	130.3
AAPL	2020	46	120.5	115.55
BRK-B	2021	32	291.81	287.01

only showing top 20 rows

func.max("Open"): Finds the highest opening price for each stock in each week (WeekHigh).

func.min("Open"): Finds the lowest opening price for each stock in each week (WeekLow).

This captures weekly price variations, which are critical for short-term trading strategies and highlights short-term trends and market behavior.

### Calculating Weekly Price Spread

```
weekly.withColumn("Spread", weekly['WeekHigh'] - weekly['WeekLow']).show()
```

```
+-----+-----+-----+-----+-----+-----+
|Ticker|Year|Week|WeekHigh|WeekLow|    Spread|
+-----+-----+-----+-----+-----+-----+
| BRK-B|2022| 14| 352.0| 341.17| 10.829987|
| BRK-B|2022| 10| 326.59| 322.49| 4.100006|
| BRK-B|2021| 14| 264.22| 260.02| 4.200012|
| META|2022| 43| 131.68| 97.98| 33.69999|
| META|2020| 6| 212.51| 203.44| 9.069992|
| TSLA|2022| 20| 255.72| 235.67| 20.050003|
| TSLA|2020| 19| 52.92| 46.73| 6.1899986|
| TSLA|2020| 16| 51.49| 39.34| 12.150002|
| TSLA|2018| 39| 20.86| 18.02| 2.8400002|
| BRK-B|2018| 48| 217.23| 209.3| 7.9299927|
| MSFT|2022| 6| 309.87| 301.25| 8.619995|
| MSFT|2021| 2| 218.47| 213.52| 4.949997|
| META|2022| 40| 140.49| 136.76| 3.730011|
| AAPL|2020| 27| 91.96| 88.31| 3.6500015|
| BRK-B|2020| 19| 180.05| 173.4| 6.650009|
| MSFT|2020| 1| 158.78| 158.32| 0.45999146|
| META|2020| 36| 298.88| 287.25| 11.630005|
| AAPL|2021| 25| 134.45| 130.3| 4.149994|
| AAPL|2020| 46| 120.5| 115.55| 4.949997|
| BRK-B|2021| 32| 291.81| 287.01| 4.799988|
+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

A new column Spread was added, representing the range of price variation for each week.

This quantifies weekly volatility, aiding in risk assessment for short-term investors and allows identification of highly volatile stocks or periods with significant price changes.

The stock data after data processing:

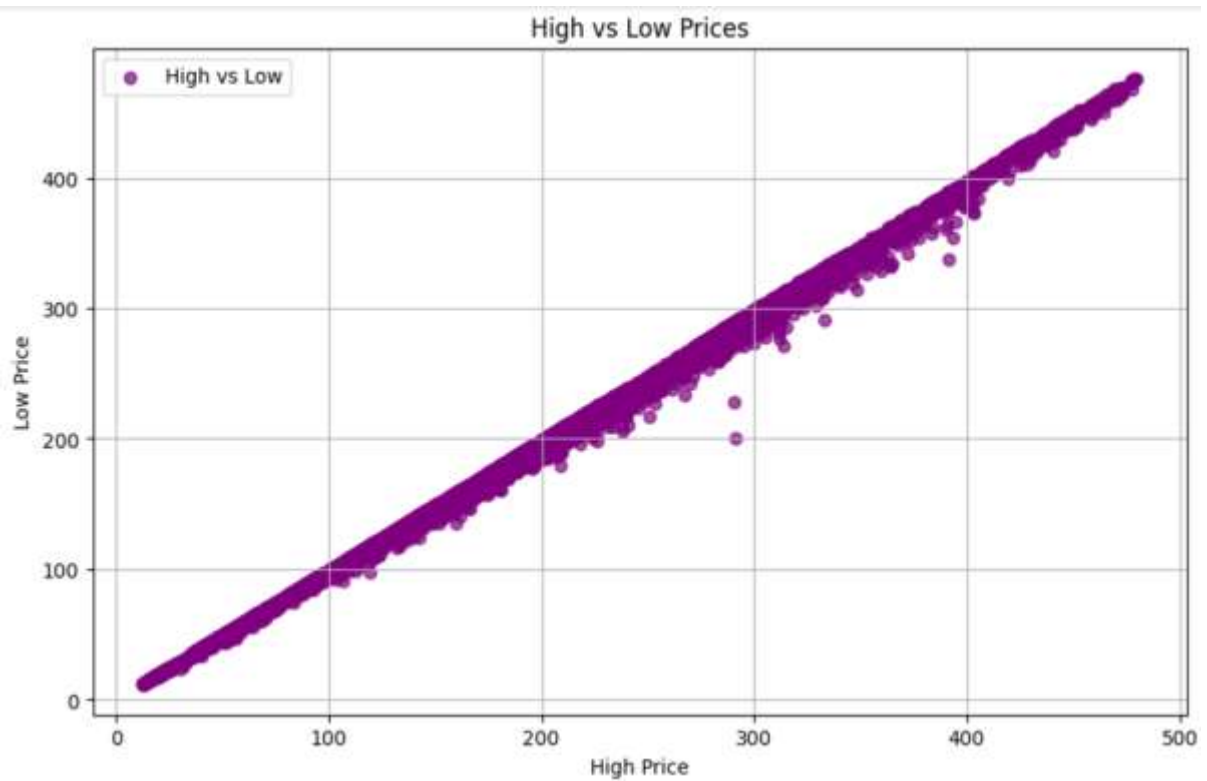
Ticker	Year	Month	Day	Week	Volume	Open	Low	High	Close	YearlyHigh	YearlyLow	WeekHigh	WeekLow	MonthHigh	MonthLow
BRK-B	2023	5	31	22	6175417	321.12	319.39	322.41	321.08	331.0	294.68	321.86	321.12	331.0	320.44
BRK-B	2023	5	30	22	3232461	321.86	319.0	322.47	322.19	331.0	294.68	321.86	321.12	331.0	320.44
BRK-B	2023	5	26	21	3229873	320.44	319.67	322.63	320.6	331.0	294.68	330.75	320.44	331.0	320.44
BRK-B	2023	5	25	21	4251935	320.56	317.71	320.56	319.02	331.0	294.68	330.75	320.44	331.0	320.44
BRK-B	2023	5	24	21	3075393	322.71	319.56	323.0	320.2	331.0	294.68	330.75	320.44	331.0	320.44

only showing top 5 rows

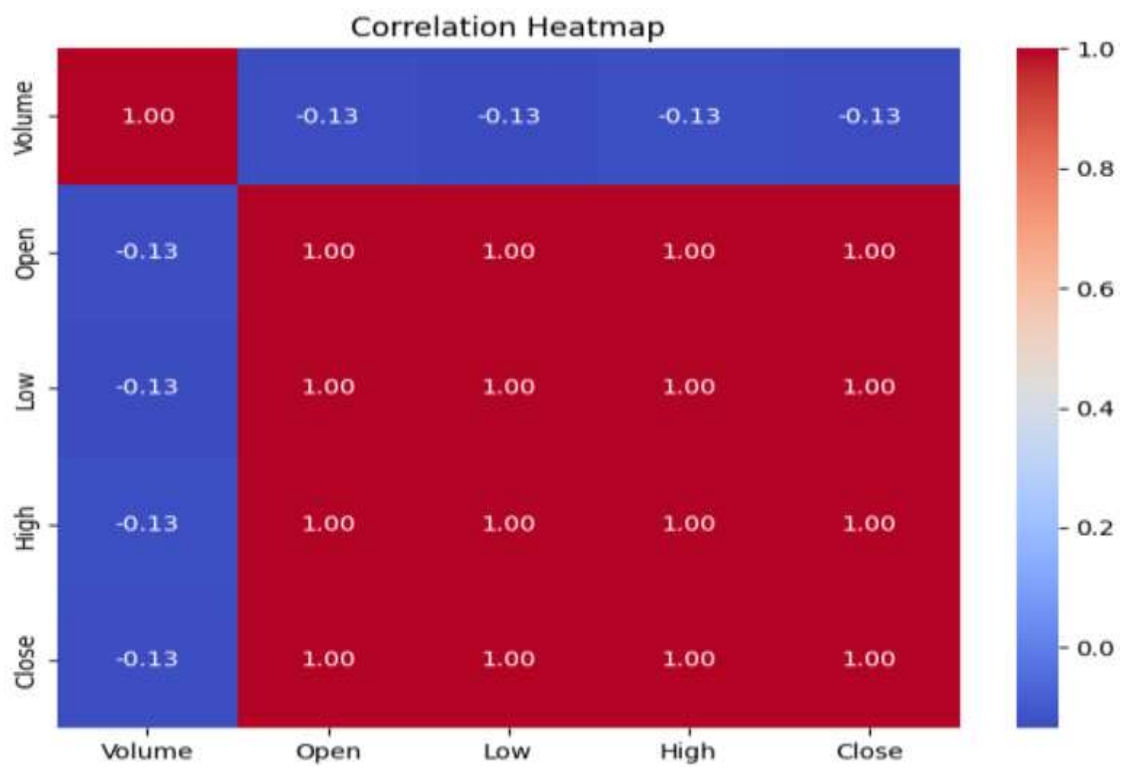
Plotting of close price vs Index:



Plotting of High vs Low Prices:




Correlation Heatmap:





The correlation heatmap identifies the correlation between the features the correlation is found for the features Volume, Close, Open, High, Low. It ranges from -1 to 1 where -1 refers to features are negatively correlated, 0 refers to no correlation, 1 refers to positive correlation.

Price Metrics over Time vs Index plotting:



By plotting a stock's price metrics (e.g., Open, Close, High, and Low) alongside the index we can identify whether the stock is outperformance or underperformance. If the stock price consistently increases at a steeper rate than the index, it indicates the stock is outperforming the market. If the stock lags the index or declines when the index rises, it suggests underperformance.

#### 4.4.4 Model Selection

##### Feature Engineering

The first step involved assembling relevant features into a single vector. The selected features were derived from exploratory data analysis (EDA)

Features Description:

Volume: Total stock volume traded.

Open, Low, High: Intraday price metrics.

YearlyHigh, YearlyLow: Maximum and minimum stock prices for the year.

WeekHigh, WeekLow: Weekly price extremes.

MonthHigh, MonthLow: Monthly price extremes.

The VectorAssembler was used to combine these features into a single column, features, for use in machine learning:



```
assembler = VectorAssembler(inputCols=feature_columns, outputCol="features")
final_stocks = assembler.transform(final_stocks)
```

## Data Preparation

The dataset was split into training and testing subsets:

```
train_data, test_data = final_stocks.randomSplit([0.8, 0.2], seed=42)
```

Training Data: 80% of the dataset, used to train the linear regression model.

Testing Data: 20% of the dataset, reserved for evaluating the model's performance.

## Model Training

A Linear Regression model was trained using the features column as the predictor and the Close column as the target variable:

```
lr = LinearRegression(featuresCol="features", labelCol="Close")
```

```
lr_model = lr.fit(train_data)
```

Linear regression is a straightforward and interpretable regression method, making it suitable for analyzing relationships between stock prices and their predictors.

A Random Forest Regressor was used for training the model:

```
rf = RandomForestRegressor(featuresCol="features", labelCol="Close", numTrees=50,
maxDepth=10, seed=42)
```

```
rf_model = rf.fit(train_data)
```

numTrees=50: The model builds 50 decision trees to make predictions, ensuring a strong ensemble approach.

maxDepth=10: Limits the depth of each tree to prevent overfitting.

seed=42: Ensures reproducibility of results.

The advantages of random forest are that it handles non-linear relationships effectively, reduces overfitting by averaging predictions across multiple trees, automatically accounts for feature importance during training.

### 4.4.5 Model Evaluation

The model was evaluated on the test data using the Root Mean Squared Error (RMSE) metric,  $R^2$ , MAE.

RMSE: Root Mean Square Error measures the average deviation between predicted and actual closing prices. Lower values indicate better model performance.

```
evaluator = RegressionEvaluator(labelCol="Close", predictionCol="prediction",
metricName="rmse")
```

```
rmse = evaluator.evaluate(predictions)
```

R<sup>2</sup> (Coefficient of Determination): This indicates how well the model explains the variance in the target variable.

```
mae_evaluator = RegressionEvaluator(labelCol="Close", predictionCol="prediction",  
metricName="mae")
```

```
mae = mae_evaluator.evaluate(predictions)
```

MAE: It represents the average absolute difference between predicted and actual values.

```
train_predictions = lr_model.transform(train_data)
```

```
train_r2 = r2_evaluator.evaluate(train_predictions)
```

```
test_r2 = r2_evaluator.evaluate(predictions)
```

#### 4.4.6 Predictions

The trained model was used to predict the closing prices for the test data:

```
predictions = lr_model.transform(test_data)
```

```
predictions.select("prediction", "Close").show ()
```

```
R2 Score for Train Data: 0.9997512142649044  
R2 Score for Test Data: 0.9998039957143493  
Root Mean Squared Error (RMSE): 1.4077940456962728  
Mean Absolute Error (MAE): 0.9190876465741273  
+-----+  
| prediction| Close|  
+-----+  
| 317.8487732700164| 317.43|  
| 321.8461201383309| 322.49|  
| 186.49289324434295| 186.02|  
| 205.6375499150663| 204.99|  
| 300.83416623086356| 299.98|  
| 312.7890788135992| 312.62|  
| 226.30252954011| 226.53|  
| 204.36388748404096| 204.1|  
| 213.3737143214504| 215.16|  
| 210.53938980510037| 209.99|  
| 226.56493042943418| 226.8|  
| 286.50767285931903| 285.31|  
| 211.715309698274| 211.8|  
| 325.4870225428558| 324.34|  
| 204.40836003143886| 203.88|  
| 305.85637176857114| 303.43|  
| 324.73671225644216| 323.22|  
| 210.51404303447995| 210.9|  
| 287.8897341763043| 288.72|  
| 302.41275580990794| 304.27|  
+-----+  
only showing top 20 rows
```

Predicted vs Actual: The predictions for closing prices were compared to the actual values to assess the model's effectiveness visually and numerically.

```
predictions = rf_model.transform(test_data)
```

```
predictions.select("prediction", "Close").show ()
```

```
R² Score for Train Data: 0.9984576295064094
R² Score for Test Data: 0.9981182439796717
Root Mean Squared Error (RMSE): 4.3620234642756825
Mean Absolute Error (MAE): 2.4587351644624267
```

```
+-----+
| prediction| Close|
+-----+
| 313.0497680522891| 317.43|
| 327.0045659653018| 322.49|
| 188.85020575874188| 186.02|
| 205.36339867059368| 204.99|
| 297.8687238727541| 299.98|
| 314.4003915447609| 312.62|
| 223.10432415227802| 226.53|
| 204.47124558801804| 204.1|
| 214.3978792744508| 215.16|
| 211.53137690487802| 209.99|
| 226.05428562293375| 226.8|
| 288.808064611046| 285.31|
| 211.48831852143542| 211.8|
| 325.3229620118367| 324.34|
| 204.80238696509107| 203.88|
| 306.5286760746083| 303.43|
| 327.0045659653018| 323.22|
| 209.5005788077691| 210.9|
| 286.19610574605383| 288.72|
| 298.1821302916428| 304.27|
+-----+
```

```
only showing top 20 rows
```

## 5 Conclusion

This project effectively showcases the integration of big data tools and machine learning techniques for comprehensive stock price analysis and prediction. The use of MapReduce allowed for efficient processing and aggregation of stock data based on monthly and yearly groupings, enabling the calculation of key metrics such as average and maximum values for Close, Volume, Open, High, and Low prices. Hive was utilized to perform insightful analyses, such as identifying the stocks with the highest average daily price range and the highest recorded prices for each year. These capabilities underline the power of big data technologies in handling large-scale datasets and generating meaningful financial insights.

The predictive modeling component, implemented using PySpark, provided actionable results. Linear Regression achieved a remarkably high  $R^2$  score of 0.9998 on test data, showcasing its ability to closely approximate the actual stock prices. Random Forest also demonstrated strong performance with an  $R^2$  score of 0.9981, though it showed a slightly higher error margin compared to Linear Regression. These outcomes highlight the utility of machine learning models in predicting stock prices with significant accuracy, offering valuable tools for investors and financial analysts to make data-driven decisions.

Overall, the project underscores the importance of combining big data analytics and machine learning to analyze and predict stock market trends effectively. By leveraging tools like MapReduce, Hive, Sqoop, and PySpark, we processed and analyzed large datasets efficiently, while machine learning models provided precise predictions. The project demonstrates the potential for further enhancements, such as incorporating additional features, exploring advanced machine learning or deep learning techniques, and implementing real-time data processing. These improvements could make the system even more robust, scalable, and applicable for real-world financial forecasting and decision-making.

## 6 References

1. Peng, Z., 2019, January. Stocks analysis and prediction using big data analytics. In 2019 international conference on intelligent transportation, big data & smart City (ICITBS) (pp. 309-312). IEEE.
2. Jose, J., Mana, S.C. and Samhitha, B.K., 2019. An efficient system to predict and analyze stock data using Hadoop techniques. International Journal of Recent Technology and Engineering (IJRTE), 8(2), pp.2277-3878.
3. Sirisha, N. and Kiran, K.V.D., 2017, December. Stock exchange analysis using Hadoop user experience (Hue). In 2017 International Conference on Intelligent Sustainable Systems (ICISS) (pp. 1141-1144). IEEE.
4. Darwiesh, A., El-Baz, A.H. and Elhoseny, M., 2024. Intelligent risk management system for enhancing performance of stock market applications. Expert Systems with Applications, 249, p.123493.